



MONTANA
STATE UNIVERSITY


Floating Point

...

Representing Decimals Efficiently


Last Lecture

- In the last lecture we looked at
 - Binary representations for *integer values*
 - We looked at hex and character representations
 - Now we are going to consider *decimal values*



Real Numbers

- Representing non-integer values
 - E.g. $\pi = 3.1415\dots$
- One obvious mechanism: dedicate a certain number of bits to the right hand side of the decimal
 - This is known as *fixed point*



Fixed Point

- Advantages
 - Simple
 - Can use the same mathematical circuitry as integers
- Disadvantages
 - Can only represent a small domain or number OR support a small number of decimal places
- Fixed point is still used for specialized computations
 - E.g. finance

0003.1415

1010.1001

Floating Point

- Realization

- What if we allowed the point to float?
- Dedicate a variable number of bits to the left hand and right hand side of the decimal point?

$$1.2345 = \underbrace{12345}_{\text{significand}} \times \underbrace{10^{-4}}_{\text{base}}^{\text{exponent}} .$$

12345,4

Floating Point

- Advantages

- Far greater range of numbers can be represented
- No wasted leading or trailing 0 bits

- Disadvantages


- More complex
- Different circuitry for mathematical operations

$$1.2345 = \underbrace{12345}_{\text{significand}} \times \underbrace{10^{-4}}_{\text{base}}^{\text{exponent}} .$$

12345,4

Floating Point History

- First known use of floating point was in an electro-mechanical computing machine designed by Leonardo Torres y Quevedo
 - Spanish engineer
 - Designed first computer game (chess) and was a pioneer in remote control
 - *How come I've never heard of him? Good question.*





Floating Point History

- Konrad Zuse
 - German engineer
 - Designed the world's first programmable computer, the Z1
 - Included 24-bit binary floating point
 - Also developed the first high-level programming language, Plankalkül
 - *How come I've never heard of him? Another good question.*




Floating Point History

- 1950s to 1980s
 - Many competing floating point standards
 - IBM 7*
 - UNIVAC
 - Etc.
- IEEE 754 - 1985
 - Established a floating point standard for the industry
 - Intel
 - Motorola



Floating Point History

- Intel 8087
 - The first x87 floating-point coprocessor for the 8086 line of microprocessors
 - 20% to 500% faster for many operations
- Eventually this functionality was folded into the main CPU with the advent of the 486 chip



Floating Point History

- William Kahan
 - Primary architect behind the IEEE 754 standard
 - The “Father of Floating Point”
 - Wrote the program *paranoia* to test floating point implementations
 - Found the floating point bug in pentium division
 - Won the Turing award for his contributions



Floating Point Today

- X86-64
 - Includes registers for floating point values
 - 128 bits (!!!)
 - XMM0-XMM7 (part of x86-32 SSE)
 - XMM8-15 (available in 64-bit mode only)




Floating Point Details

- IEEE 754 Floating Point Representation
 - Various levels of precision
 - Half - 16 bits
 - Single - 32 bits
 - Double - 64 bits
 - Extended - 80 bits
 - Quad - 128 bits

$$1.2345 = \underbrace{12345}_{\text{significand}} \times \underbrace{10^{-4}}_{\text{base}}^{\text{exponent}} .$$


Floating Point Details

- IEEE 754 Half
 - A total of 16 bits
 - 1 *sign bit*
 - 5 *exponent bits*
 - 10 *significand* (fraction) bits
- Sign bit is obvious
 - 0 positive
 - 1 negative




Exponent

- Exponent is a 5 bit value, giving $2^5 = 32$ different possibilities
 - Because the exponent is unsigned and we wish to express negative values it is *biased* at 15
 - The unsigned value has 15 subtracted from it to get the actual exponent value
 - This allows values of 15 to -14 for the exponent




Exponent

- IEEE 754 Half
 - What about exponent value 00000?
 - Typically means 0
 - Can also mean subnormal numbers
 - Very small numbers below the normal floating point range
 - What about the exponent value 11111?
 - Means infinity or NaN, depending on significand



Significand (Fraction)

- Fraction is 10 bits
- Gives us 2^{10} (1024) possible values
 - Values are expressed in terms of $x/1024$
 - E.g 0000000001 \rightarrow 1/1024th
- Significand value is added to 1 to get a number somewhere between 1 and 2
 - Except subnormal case, when it 0 is added to it



Float Value Calculation

- Consider this 16 bit floating point number
- Sign bit: 0 (positive)
- Exponent 1 $\rightarrow 2^{(1-15)} = 2^{-14}$
- Fraction = 0000000000 $\rightarrow (1 + 0/1024)$

0 00001 000000000000₂

$$2^{-14} \times \left(1 + \frac{0}{1024}\right)$$

0.000061035156

Binary Fractions

- We are used to decimal notation
- What does 1.2345 mean in terms of fractions?
 - $1 + 2,245/10,000$
- What does 1.01011 mean in binary?
 - $1 + (01011/100000)$ *binary*
 - $1 + (11/32)$ *decimal*
 - $1 + .34375$ *decimal*
 - 1.34375


0 00001 000000000000₂

$$2^{-14} \times \left(1 + \frac{0}{1024}\right)$$

0.000061035156

Binary Fractions

- Another example:
- 0.1011
 - 1011/10000 binary
 - $11/16 = 0.6875$
- Or look at it in terms of the twos places...



Float Value Calculation

- Consider another 16 bit floating point number
- Sign bit: 0 (positive)
- Exponent 13 $\rightarrow 2^{(13-15)} = 2^{-2}$
- Fraction = 0101010101 $\rightarrow (1 + 341/1024)$
- *This is the closest 16 bit floating point can represent 1/3rd*

0 01101 0101010101₂

$$2^{-2} \times \left(1 + \frac{341}{1024}\right)$$

0.33325195

Float Precision

- This is a serious limitation of floating point: It can only be an approximation of many values

0 01101 0101010101₂

$$2^{-2} \times \left(1 + \frac{341}{1024}\right)$$

0.33325195

Float Rounding Rules

- Floating Point has different rounding rules
 - Round to nearest, ties to even
 - Round to nearest, ties away from zero
 - Round toward 0
 - Round toward $+\infty$
 - Round toward $-\infty$

0 01101 0101010101₂

$$2^{-2} \times \left(1 + \frac{341}{1024}\right)$$

0.33325195

Float Precision Implications

- You must be very careful when using floating point!
- Floating point is a bad idea when dealing with fixed precision numbers
 - E.g. Money!
- That's OK though, our industry is smart enough to understand this...

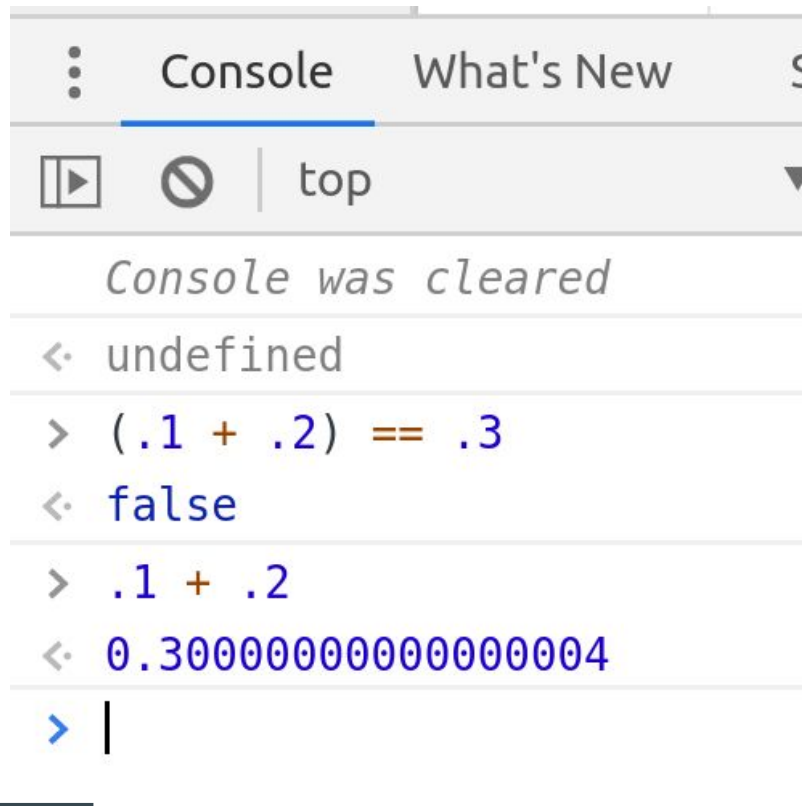
0 01101 0101010101₂

$$2^{-2} \times \left(1 + \frac{341}{1024}\right)$$

0.33325195

Float Precision Implications

- You must be very careful when using floating point!
- Floating point is a bad idea when dealing with fixed precision numbers
 - E.g. Money!
- That's OK though, our industry is smart enough to understand this...

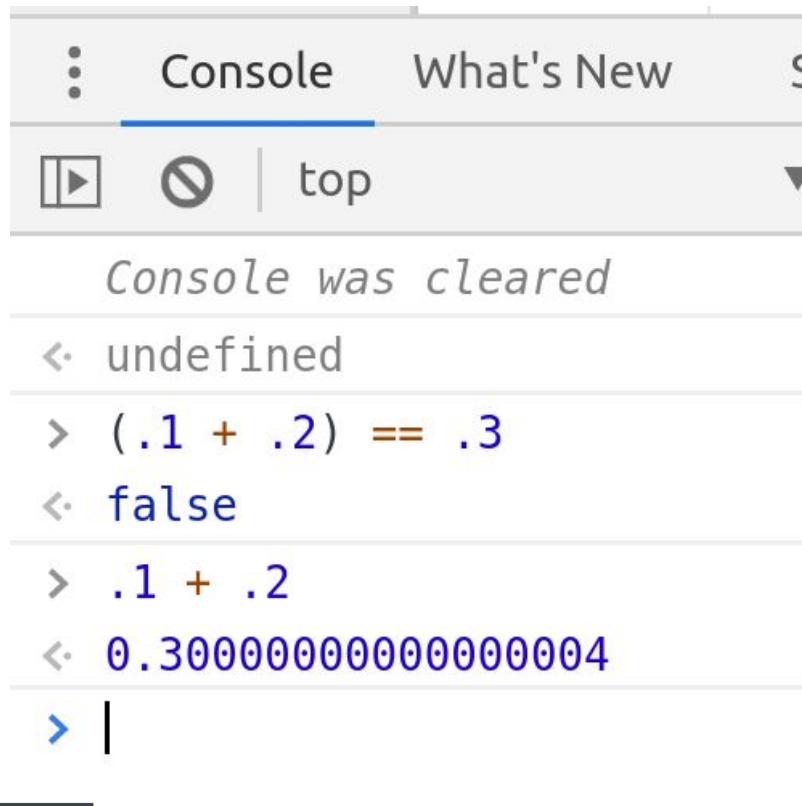


```
⋮ Console What's New S
⏮ ⓧ | top
    Console was cleared
< undefined
> (.1 + .2) == .3
< false
> .1 + .2
< 0.30000000000000004
> |
```

The screenshot shows a web browser's developer console. At the top, there are tabs for 'Console' (selected) and 'What's New'. Below the tabs, there are icons for a play button, a disabled icon (a circle with a slash), and a 'top' link. The console output shows a message 'Console was cleared' in italics. Below that, there are several lines of code and their corresponding outputs. The first line is '< undefined'. The second line is '> (.1 + .2) == .3' followed by '< false'. The third line is '> .1 + .2' followed by '< 0.30000000000000004'. The final line is '> |' with a cursor.

Float Precision Implications

- Uhhhhh
- Javascript uses 32 bit floats for numbers
- So *obviously true* mathematical statements are... false
- People are increasingly writing software in javascript...



```
⋮ Console What's New S
⏮ ⛔ top
Console was cleared
< undefined
> (.1 + .2) == .3
< false
> .1 + .2
< 0.30000000000000004
> |
```

The screenshot shows a web browser's developer console. The 'Console' tab is active. The console has been cleared, as indicated by the message 'Console was cleared'. The first log entry is 'undefined'. The second log entry shows the result of the expression `(.1 + .2) == .3`, which is `false`. The third log entry shows the result of the expression `.1 + .2`, which is `0.30000000000000004`. The console is currently empty, with a prompt character `>` and a cursor.

Float Precision Implications

> mfw



Float Precision Implications

- The Patriot Missile Incident
 - February 1991
 - Precision issue in a MIM-104 Patriot missile battery prevented it from intercepting an incoming SCUD missile
 - 28 soldiers killed



Float Special Values

- Float has some *special* values
 - A signed 0 value
 - Equal to one another
 - Division by one or the other leads to a signed infinity
 - Infinities
 - + and - infinities
 - Satisfy standard infinity math (e.g $3 + +\infty = +\infty$)
 - NaN (Not a Number)
 - Represents invalid values such as $1/0$ or $\text{sqrt}(-1)$


0 01101 0101010101₂

$$2^{-2} \times \left(1 + \frac{341}{1024}\right)$$

0.33325195


Larger Floats

- We have been looking at half floats, but this same logic generalizes to any length of bits
 - float - typically 32 bits
 - double - 64 bits
- With more bits, more precision




Larger Floats

- We have been looking at half floats, but this same logic generalizes to any length of bits
 - float - typically 32 bits
 - double - 64 bits
- With more bit, more precision
 - Still not perfect precision however



Floating Point Today

- Recall, on X86-64
 - Floating point values are passed in separate registers
 - 128 bits of precision are available
 - XMM0-XMM7 (part of x86-32 SSE)
 - XMM8-15 (available in 64-bit mode only)
 - Separate assembly instructions for working with them



Floating Point

- We took a look at how to represent non-integer values in binary
- Initially *Fixed Point* was used
- *Floating Point* representation is more efficient
 - But also more complex!
- We took a look at 16 bit floats
 - Larger floats are just more of the same
- And we looked at problems with Floating Point precision
 - Javascript is a very terrible programming language
- *REMEMBER: IT'S JUST BITS*



MONTANA
STATE UNIVERSITY